

Character Device Drivers: An Introduction



ECE 373

Prelims

- Questions on homework or reading assignments?
- Questions on class?



Char drivers

- Driver for char devices
- Typical types of char drivers
 - Mice
 - Keyboards
 - Printers
 - Memory
- Special files in /dev marked with notation:



```
- crw-rw-rw- 1 root wheel 5, 49 Apr 13 00:39 ptys1
```

Char drivers

- Driver for char devices
- Typical types of char drivers
 - Mice
 - Keyboards
 - Printers
 - Memory
- Special files in /dev marked with notation:



```
- crw-rw-rw- 1 root wheel 5, 49 Apr 13 00:39 ptys1
```

Device nodes, deeper

- Major and minor numbers used to identify device

- `crw-rw-rw- 1 root wheel 5, 49 Apr 13 00:39 ptys1`

- Major = module (tty, clock, serial ports, disk)
 - Minor = which specific instance (pty49)

- Kernel data type to store dev node is "dev_t"
- Defined in `include/linux/types.h`
- Unsigned 32-bit number, packed
 - 20-bits of minor, 12-bits of major
- Current list in `/proc/devices`

Manipulating dev nodes

- Macros used to extract numbers from `dev_t`
 - `MAJOR(dev_t dev)`
 - `MINOR(dev_t dev)`
- Found in `include/linux/kdev_t.h`
- Ensures device representation is portable
- Example 1!

Registering a new char device

- First, allocate a region

```
- int register_chrdev_region(dev_t first,  
                             unsigned int count,  
                             char *name)
```

- Uses predetermined device nodes
- Not the right way to do most things anymore...
- Requires mknod to be used
 - man mknod
- Example 2!

Properly registering char device

- Request a region to be allocated for you

```
- int alloc_chrdev_region(dev_t *dev,  
                           unsigned int firstminor,  
                           unsigned int count,  
                           char *name)
```

- Auto-fills "dev" with device nodes
- "firstminor" typically 0, can be anything

Cleaning up after your char device

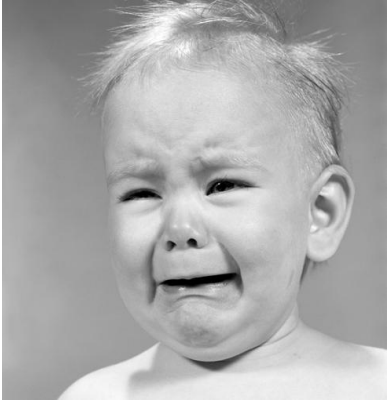
- With any dynamic stuff, it must be cleaned up
 - `void unregister_chrdev_region(dev_t first, unsigned int count)`
- Destroys internal kernel references
- Resource "leak" if not called
- Example 3!



It's there, I want to use it!

- Well, module is pretty dumb right now...
- Device needs to be connected
- Use mknod to create device `'/dev/ece'`
 - `man mknod...`
- Can you read and write it?





Why didn't it work?

- `alloc_chrdev_region()` just configures device internal to kernel
- No linkage to anything in the upper device subsystem
- System calls and driver callbacks

Beginning to hook it all up

- Structure "`file_operations`" provides function pointers into system call interface
- Main linkage into `/dev` filesystem for char drivers
- Driver does not need to implement all of them
- Behaves similarly to object-oriented code




Snippet of file_operations API

```
struct file_operations {  
    struct module *owner;  
  
    int (*open) (struct inode *, struct file *);  
  
    int (*release) (struct inode *, struct file *);  
  
    ssize_t (*read) (struct file *,  
                     char __user *,  
                     size_t, loff_t *);  
  
    ssize_t (*write) (struct file *,  
                      const char __user *,  
                      size_t, loff_t *);  
  
}
```

Files and inodes

- Internal kernel structures to manage "files"
- Opened files are managed by "`struct file`" internal to kernel
- Inode is the describer of a physical file on disk
- Many directory entries can reference single inode, inode cannot refer to more than one file



The cdev!

- Yes, another structure to worry about
- Struct that represents char devices inside kernel
- Initialized with:
 - `void cdev_init(struct cdev *cdev,
 struct file_operations *fops)`
- Added with:
 - `int cdev_add(struct cdev *cdev, dev_t num,
 unsigned int count)`
- Cleaned up with:
 - `void cdev_del(struct cdev *cdev)`

Last minute safety checks

- `file_operations` must be configured and ready to go before `cdev_add()`
- Could run into NULL pointer exceptions
- Kernel will go boom if this is misconfigured
- Example 4!



Passing data through

- System call only passes data
- Need mechanism to copy data into kernel buffers, out of kernel buffers
- Two handy-dandy functions:
 - `copy_from_user`
 - `copy_to_user`
- Example 5!



More info

- Linux Device Drivers, 3rd Edition
 - Chapter 3, Pages 42 – middle of page 57
- Essential Linux Device Drivers
 - Chapter 5, Pages 119 – top of page 129

Coming soon...

- Hardware access capabilities
- I/O ports, MMIO, PCI
- Readings:
 - LDDR Chapter 9, Chapter 12
 - ELDD Chapter 10

