

Threads and Concurrency



ECE 373

Threads of Execution

- Thread – smallest unit of processing that can be scheduled by the OS
 - UNIX-like systems are process oriented
 - WinNT-like systems are thread oriented
- Single CPU core
 - single line of instructions at any one time
- Multiple CPU cores
 - multiple simultaneous threads of execution



Thread Sources

- Kernel threads
- User processes
- Workqueue threads
- Timer callbacks
- Interrupt handlers



Threads




- Kernel threads
 - Jobs the kernel itself is doing for internal [projects]
 - All have access to the same kernel data
- User context
 - Threads running user jobs, might be running kernel code to service system calls from user code
 - Collected into specific user processes, see only the individual process space data
- Interrupts
 - Not really full threads, but in the mix (unless threaded interrupts...)
- See `ps -ef`

Linux process tree




Thread simple



OS Scheduler

- Chooses which to process/thread to run next on which CPU


- Ready queue
- Wait queue
- Thread priority
- Time slice
- Preemption
- CPU core affinity
- Etc



OS Scheduler


- Chooses which to process/thread to run next on which CPU

- Ready queue
- Wait queue
- Thread priority
- Time slice
- Preemption
- CPU core affinity
- Etc




- Interrupts mess with scheduler plans

Thread, interrupted




Concurrency and conflict

- Multiple threads could hit data at same time



Concurrency and conflict

- Multiple CPU threads could hit data at same time



Example: conflict

- User thread 1 asks for data
 - Driver requests data from HW, sleeps while waiting
- User thread 2 removes device
 - Driver removes data structures
- HW interrupt to finish data retrieval
 - Driver interrupt handler tries to access removed data struct
 - Uh oh...

Order matters

Instance 1	Instance 2	Value
read very_important_count		5
add 5 + 1 = 6		6
write very_important_count		6
	read very_important_count	6
	add 6 + 1 = 7	7
	write very_important_count	7

Order matters

Instance 1	Instance 2	Value
read very_important_count		5
add 5 + 1 = 6		6
write very_important_count		6
	read very_important_count	6
	add 6 + 1 = 7	7
	write very_important_count	7

Instance 1	Instance 2	Value
read very_important_count		5
	read very_important_count	5
add 5 + 1 = 6		6
	add 5 + 1 = 6	6
write very_important_count		6
	write very_important_count	6

Order matters

Instance 1	Instance 2	Value
read very_important_count		5
add 5 + 1 = 6		6
write very_important_count		6
	read very_important_count	6
	Add 6 + 1 = 7	7
	write very_important_count	7

Instance 1	Instance 2	Value
read very_important_count		5
add 5 + 1 = 6		6
write very_important_count		6
	read very_important_count	5
	Add 5 + 1 = 6	6
	write very_important_count	6




Coordination needed

- Multiple threads could hit data at same time
- Tools:
 - Completions
 - Semaphore
 - Atomic action – increment, decrement
 - Mutex
 - Spin lock
 - RCU



Waiting for the completion



Atomic action

- CPU instructions for atomic increment, decrement, test_and_set
 - <https://elixir.bootlin.com/linux/latest/source/arch/x86/include/asm/atomic.h#L95>
- All cores must coordinate CPU cache – expensive

```
atomic_inc(x)
atomic_inc_and_test(x)
set_bit(n, *s)
clear_bit(n, *s)
test_bit(n, *s)
```

```
static inline void atomic_inc(atomic_t *v)
{
    asm volatile(LOCK_PREFIX "incl %0"
                 : "+m" (v->counter));
}
```

Semaphore



Semaphore



- Counter that many threads can inc/dec
 - Usually starts positive, each user decrements to start, inc when done
 - If zero, next process must wait
- Thread A might start operation, sleep, then thread B might finish
- Use atomic inc/dec to implement counter

```
struct semaphore sem;  
sema_init(&sem, val);  
down(&sem);  
up(&sem);
```

Spinlock

Spinlock

- While not have lock, try again
 - Tight spin
 - Unlimited spin can "hang" thread, block other operations

```
while (test_and_set(3, &bit_string))  
    /* tight loop */ ;
```

- Alternative is a sleep spin

- Less CPU intense...

```
while (test_and_set(3, &bit_string))  
    usleep(2);
```

- ... but might miss a window of opportunity

- Linux:

```
spinlock_t slock;  
spin_lock_init(&slock);  
spin_lock(&slock);  
spin_unlock(&slock);  
spin_trylock(&slock);
```

Mutex

Mutex

- Mutual Exclusion
 - Everyone waits until the thread is done
 - Thread A gets lock, only thread A can release it
- Other threads will sleep while waiting for lock
- Good for blocking access to data, other resource
- Like spinlock, but more restrictive

- Linux:


```
struct mutex mlock;  
mutex_init(&mlock);  
mutex_lock(&mlock);  
mutex_unlock(&mlock);
```


Deadlock

Deadlock

- Deadlock possible with multiple locks
 - Process A gets lock 1, wants lock 2
 - Process B gets lock 2, wants lock 1

- Now what?
 - Linux does some checking



Watch out



- Blocking processes
 - Slowing other threads, whole system
 - Priority inversion – low prio thread holds lock, high prio thread can't continue
- Granularity
 - Lock smallest amount of code possible
- Balancing lock/unlock
- CPU communication overhead
- Hard to debug because of timing related

Reading

- LDD3 – Chapter 5: Concurrency
- LDD3 – Chapter 10: Interrupts
- ELDD – Chapter 2: Concurrency, pgs 39-48
- ELDD – Chapter 3: Kernel Facilities
- ELDD – Chapter 4: Interrupt Handling, pgs 92-103
- Linux src – ../Documentation/atomic_ops.txt

