

Final Review



ECE 373

Prelims

- Questions?
- The end is near!!!



What's fair game

- The entire term
- Focus more on material after mid-term
- Be prepared for anything though
- Open book, open notes ✕
- Touching on high points today

Bits and Pieces

- Minimal callback hooks and compile headers

- Basic #includes *not really on exam — Art of wild pr*
- MODULE_LICENSE(lic) – legal strings *read with mod info*
- *rmmod* __init, __exit *rmmod*
- module_init(func), module_exit(func)


- Time values ***

- Jiffies, HZ *Jiffies / Sec. , 1000 Jif/sec = 1000 Hz. 1000 Jif/sec = 1000 Hz. 1000 Jif/sec = 1000 Hz.*
Kernel wall ind. tick — time rep. inside running kernel
100-250 — server
s.w. int. fire in longer intervals

KNOW CONVERSION

Main Bus -

PCI



Config ID block
(SUB) Vendor { Id's
device

- Communication/connection method for devices
- Devices use both port mapped and memory mapped I/O (focus on MMIO only for final)
- *get this* BAR – base address register
 - Starting address for device memory map
 - *then* Driver uses `writel()` and `readl()` to access device registers with MMIO
- Callbacks
 - *pci side - implemented in driver* Probe, remove

Kernel

Concurrency

~ 5:26

Order matters

t
↓

Instance 1	Instance 2	Value
read very_important_count		5
add 5 + 1 = 6		6
write very_important_count		6
	read very_important_count	6
	Add 6 + 1 = 7	7
	write very_important_count	7

RACE
CONDITION?

Instance 1	Instance 2	Value
read very_important_count		5
add 5 + 1 = 6		6
write very_important_count		6
	read very_important_count	5
	Add 5 + 1 = 6	6
	write very_important_count	6



Atomic action

- CPU instructions for atomic increment, decrement, test_and_set
- All cores must coordinate CPU cache – expensive

- `atomic_inc(x)`
- `atomic_inc_and_test(x)`
- `set_bit(n, *s)`
- `clear_bit(n, *s)`
- `test_bit(n, *s)`

memory fences/barriers
BML

—— manip. of data gets synchronized

not safe \rightarrow allows thread to sleep. not in interrupt

Locks and synchronization

- Mutex's
- Spin locks
- Semaphores
- Completions
- RCU — *loading type*
- When is it safe to lock?

Accessed
data modified only one
at a time
protects code & shared
the data (critical sections)

cpu burns,
spins @ 100%
until lock is
achieved

Delayers

- `mdelay()` , `udelay()` , `ndelay()`
- While loop on a counter, no scheduler action
 - <http://lxr.linux.no/#linux+v2.6.38/include/linux/delay.h#L46>
 - <http://lxr.linux.no/#linux+v2.6.38/arch/x86/lib/delay.c#L116>
- Only way to get short period delays
- Not very friendly for long periods
- Could block jiffies update if interrupts disabled

Sleepers

→ Kernel, ms. $1s = 1000!$

- `msleep()`, `usleep()`
- Loop that can be scheduled out
- Friendly for long periods
- Not good to use in interrupt context (both SOFTIRQ and HARDIRQ contexts)

WARN, BUG

- Code warnings
 - `BUG()`, `BUG_ON(expr)`
 - `WARN()`, `WARN_ON(expr)`, `WARN_ONCE()`
- BUG stops kernel thread
- Both produce stack dump output
- WARN does not stop kernel thread

vmalloc
↳ Virt. Ker.
kmalloc

DMA in Linux

- Function APIs exist for each driver type to control DMA
- DMA consists of mapping memory for DMA, unmapping when finished
- Mapping memory means pinning it down, not allowing it to be swapped out by memory manager
- DMA deals with physical address (or bus address)

~ 5:37

What is a userspace driver?

- A way to drive hardware from outside kernel
- Accesses various resources
 - I/O ports
 - Memory regions
 - Control interfaces
- Resources presented by kernel scanning hardware in standardized ways
 - e.g. PCI bus scan
- Not a .ko file!

Mapping a file is as simple as...

- `open()` the file

** mmap()* the file *no read/write (on kernel)*

- Inspect the memory region returned from `mmap()`
- `munmap()` the file
- `close()` the file

Why use a userspace driver?

- Quick prototyping
- Free from kernel ABI changes
 - Application Binary Interface
- Typically doesn't blow up the machine
 - Constrained to mapped space
- Ease of use
 - Can be written in most compiled and scripted languages
- May not have ability to change the kernel

Adding
new
drivers

Types of OS's in the wild

- Single-user (Phones, PC's)
- Multi-user (Servers, mainframes, "cloud")
- Real-time (Stop lights, shuttle navigation)
- Embedded (Watch, routers, car engine, mp3)



HW Interrupts

- Hardware wants attention
 - Data waiting, might be time-sensitive
- Interrupt handlers
 - Temporarily take over current thread, whether kernel or user
 - Can't be scheduled, can't sleep



Interrupts: Be quick!

- Blocking other interrupt handling and user job
- Grab HW info, stash away for later
- Wake up driver code with worker thread or waiting on a lock
- Don't call code that might sleep
 - `*sleep()`, `kmalloc()`, other I/O functions
- Locks?
 - Be careful...



Character drivers

- Typical types of char devices
 - Mice
 - Keyboards
 - Printers
- Stream data to and from device, no set size
- Links file_operations through struct cdev

Beginning to hook it all up

- Structure `"file_operations"` provides function pointers into system call interface
- Main linkage into `/dev` filesystem for char drivers
- Driver does not need to implement all of them
- Behaves similarly to object-oriented code



Snippet of file_operations

```
. struct file_operations {  
    - struct module *owner;  
  
    - int (*open) (struct inode *, struct file *);  
  
    - int (*release) (struct inode *, struct file *);  
  
    - ssize_t (*read) (struct file *,  
                      char __user *,  
                      size_t, loff_t *);  
  
    - ssize_t (*write) (struct file *,  
                      const char __user *,  
                      size_t, loff_t *);  
  
. }
```

Block drivers

- Drivers that transfer fixed-block sizes
- Primarily for disk and storage devices
- Block I/O is mostly because of how disks are laid out
- Spindle drives would write in small clusters
- Clusters create blocks
- Not as necessary on modern drives



The Kconfig framework

- Complex infrastructure to enable/disable kernel features
- Used to manipulate makefiles
- Layered, like an onion (and stinky too!)
- Can implement multiple dependencies

Building and booting a kernel

- make, make modules_install, make install
- GRUB
- vmlinuz and vmlinux images
- Modules installation
- Initial RAM disks (initrd)

Add your own code!

- How to add new pieces to the kernel
- Editing/adding Kconfig
- Creating your makefile
- Enjoying your time with maintainers...

Descriptors

- What is a descriptor?
 - Hardware field describing what work to do
 - Hardware field describing what work was done
- Carries bits and fields
- Carries pointers to buffers needing to be DMA'd into hardware
- Carries pointers to buffers DMA'd out of hardware

Filling it out

```
do {
    buffer_info = &tx_ring->buffer_info[i];
    tx_desc = E1000_TX_DESC(*tx_ring, i);
    tx_desc->buffer_addr = cpu_to_le64(buffer_info->dma);
    tx_desc->lower.data = cpu_to_le32(txd_lower |
                                     buffer_info->length);
    tx_desc->upper.data = cpu_to_le32(txd_upper);

    i++;
    if (i == tx_ring->count)
        i = 0;
} while (--count > 0);

tx_desc->lower.data |= cpu_to_le32(adapter->txd_cmd);
```

That should do it...

